

**Introduction to iCER  
and the HPCC**

John B. Johnston      Dr. Dirk Colbry  
HPC Programmer      Research Specialist  
Institute for Cyber-Enabled Research

MICHIGAN STATE UNIVERSITY

© 2010 Michigan State University Board of Trustees.

iCER

## Agenda

- Introduction to HPCC/iCER
- Bioinformatics Support at iCER
- Doing more faster - Pleasantly Parallel

MICHIGAN STATE UNIVERSITY

iCER

## Science and Computation

- Pillars of Science:
  - Theory
  - Experimentation
  - Simulation
  - (Big) Data

The quest for knowledge used to begin with grand theories. Now it begins with massive amounts of data. Welcome to the Petabyte Age.  
Wired Magazine, 2009

MICHIGAN STATE UNIVERSITY

iCER

## What is Advanced Computing Hardware?

- Anything more advanced than your desktop
- Local resources
  - Lab, Department, Institution (HPCC)
- National resources
  - NSF (XSEDE, Blue Waters), DOE (Jaguar) , Others
- Commercial Resources (cloud computing)
  - Amazon, Azure, Liquid Web, Others

MICHIGAN STATE UNIVERSITY

iCER

## Why use Advanced Computing Hardware?

- Science takes too long
- Computation runs out of memory
- Need licensed software
- Need advanced interface (visualization)

MICHIGAN STATE UNIVERSITY

iCER

## Types of problems

- CPU bound
  - Lots of computing (think simulation)
- Memory bound
  - Requires lots of memory (think genomics)
- I/O bound
  - Requires lots of data (think astronomy)

\*\* Many problems fall in more than one category

MICHIGAN STATE UNIVERSITY

iCER

## Types of Systems

**Local Resources**

- Special Use Systems
  - SMP – Symmetric Multiprocessor (fat node)
  - GPGPU (General Purpose Graphics Processing Unit)
- Typical HPC cluster
  - Commodity computers
  - High speed backbone
  - High speed network storage

**National and Commercial**

- Advanced HPC
  - Specialty hardware
  - High speed backbone
  - High speed storage
- Grid
  - Many HPC systems linked together by high speed network
- Cloud
  - Lots of definitions
  - Typically refers to computing as a service using highly flexible virtual machines



## 1957 MISTIC Mainframe

- MSU's first mainframe
- Hand built by grad students
  - Dick Reid
  - Glen Keeney




## After MISTIC

- 1957 MISTIC
- 1963-1973 CDC 3600
- 1967 Computer Science Department
- 1968 CDC 6500
- 1971 MERIT
- 1978 Cyber 750
- 2004 HPCC**
- 2009 ICER**



## 2004 MSU HPCC

- Provide a level of performance beyond what you could get and reasonably maintain as a small group.
- Provide a variety of technology, hardware and software, that would allow for innovation not easily found.

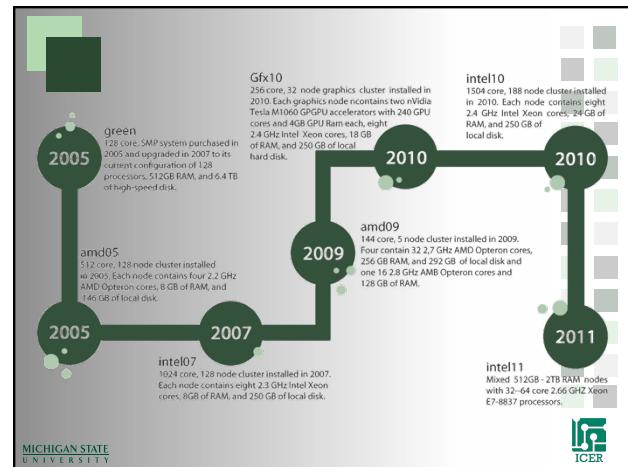
10

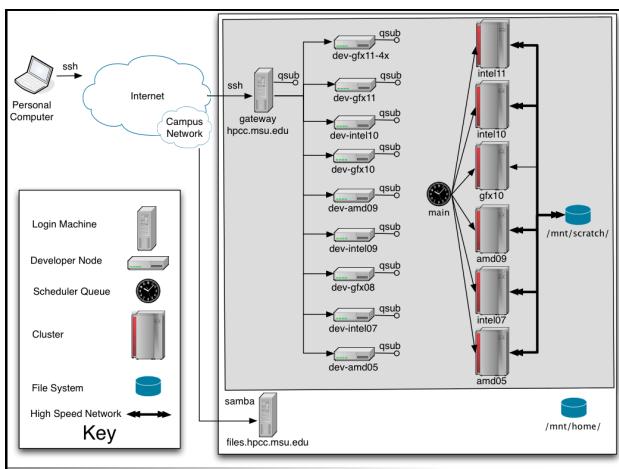


## HPCC Hardware

**FREE\***

- Over 3110 Nodes
- Over 23072 CPU cores
- Over 33 TB of scratch memory space
- GPGPU cluster with 64 Tesla Nodes
- 2 TB Shared memory machines
- 50GB backed up home directory space



## 2009 iCER

The Institute for Cyber Enabled Research (iCER) at Michigan State University (MSU) was established to coordinate and support multidisciplinary resource for computation and computational sciences. The Center's goal is to enhance MSU's national and international presence and competitive edge in disciplines and research thrusts that rely on advanced computing.

**MICHIGAN STATE UNIVERSITY**

**iCER**

## Bigger Science



- The goal of iCER is NOT:
  - Kflops
- Instead, the goal of iCER IS:
  - KSciences / second
- Doing More Science, Faster
  - Reducing the “Mean time to Science”
- iCER is designed to help researchers do their science and when appropriate scale them up to one of the national labs

**MICHIGAN STATE UNIVERSITY**

**iCER**

## Available Software

- Center Supported Development Software
  - Intel compilers, openmp, openmpi, mvapich, totalview, mkl, pathscale, gnu...
- Center Supported Research Software
  - MATLAB, R, fluent, abaqus, HEEDS, amber, blast, ls-dyna, starp...
- Customer Software
  - gromacs, cmake, cuda, imagemagick, java, openmm, siesta...

– For a more up to date list, see the documentation wiki:
 

- <http://wiki.hpcc.msu.edu/>

**MICHIGAN STATE UNIVERSITY**

**iCER**

## Accounts

- Each user has 50Gigs of backed-up personal hard drive space.
  - /mnt/home/username/
- Users have access to more than 33TB of high speed parallel scratch space.
  - /mnt/scratch/username/
- Shared group space is also available upon request.

**MICHIGAN STATE UNIVERSITY**

**iCER**

## Buy-In Opportunities

- Every 6 months to a year
- Researchers get exclusive use of their nodes within 4 hours of submitting a job
- Buy-in jobs will automatically overflow into the general resources.

**MICHIGAN STATE UNIVERSITY**

**iCER**

**Beyond Local Resources**

- National resources
  - NSF (XSEDE, Blue Waters)
  - DOE (Jaguar)
  - Others
- Commercial Resources (cloud computing)
  - Amazon
  - Azure
  - Liquid web
  - Others

MICHIGAN STATE UNIVERSITY

 iCER

**More than Just Hardware**

- Education
- Grant Writing / Support
- Advancing Computational Science

MICHIGAN STATE UNIVERSITY

 iCER

**VSCSE**



**VIRTUAL SCHOOL  
OF COMPUTATIONAL  
SCIENCE AND  
ENGINEERING**



- **Science Cloud:** July 30-August 3
- **Many-core Programming:** August 13-17

MICHIGAN STATE UNIVERSITY

 iCER

**CI-Days, Save the Date!!!**

- October 25-26
  - CI Days is an opportunity for the MSU research community to share information about cyberinfrastructure resources and their use in research.
  - All events will be held on campus in the Biomedical and Physical Sciences (BPS) building.

<http://tech.msu.edu/CI-Days/>

MICHIGAN STATE UNIVERSITY

 iCER

**People**

- iCER Director
  - Dr. Wolfgang Bauer
- HPCC Director
  - Dr. Bill Punch
- Administrative Assistant
  - Kelly Osborn

MICHIGAN STATE UNIVERSITY


 iCER

**HPC Administrators**

- Andrew Keen
- Jim Leikert
- Gregory Mason

MICHIGAN STATE UNIVERSITY


 iCER

**iCER Research Specialist**

Ben Ong and Me

- Research Consulting
- HPCC Programming
- Proposal Writing
- Training and Education
- Outreach




**MICHIGAN STATE UNIVERSITY**

**iCER**

**HPC Programmers**

- John Johnson
  - 50% Biology
  - 50% iCER
  - Primary liaison to Biology



**MICHIGAN STATE UNIVERSITY**

**iCER**

**iCER Bioinformatics Support**

**MICHIGAN STATE UNIVERSITY**

© 2012 Michigan State University Board of Trustees.

**iCER**

**Bioinformatics Outreach**

- Biological analysis software
- Common data set repositories
- Database services
- Virtual Machines
- Research-based storage space
- Seminars
- One-on-one Consulting
- Programming and scripting assistance

[wiki.hpc.msu.edu/display/Bioinfo/Bioinformatics+Support+at+MSU](http://wiki.hpc.msu.edu/display/Bioinfo/Bioinformatics+Support+at+MSU)

**MICHIGAN STATE UNIVERSITY**

**iCER**

**Some Examples**

- Multithreaded BLAST – shared memory
- mpiBLAST – distributed memory
- Velvet Assembler – multithreaded shared mem
- MAKER2 – MPI, distributed memory
- Parallel Solutions: OpenMP, OpenMPI, MVAPICH, MVAPICH2

**MICHIGAN STATE UNIVERSITY**

**iCER**

**Virtual Machines**

- Virtual “Servers” expressed in software
- Available for research labs/working groups
- Flavors currently available:
  - Galaxy (single node only)
  - BLAST (web browser based)
  - UCSC/GBrowse Genome Browsers
  - MySQL database services (postgres can be provided also)

**MICHIGAN STATE UNIVERSITY**

**iCER**

## Tutorials

- Titus Brown's ANGUS-NGS tutorials, converted for using examples on HPC instead of Amazon
- Using UCSC for certain tasks
- mpiBLAST
- Velvet and Oases
- Others being developed...

MICHIGAN STATE UNIVERSITY

ICER

## Bioinformatics Contact

- John Johnston, HPC Programmer
  - M-W, 1449 BPS, 884-2572
  - Th-F, 505 BMB, 432-7177
  - [johnj@msu.edu](mailto:johnj@msu.edu)
- Ticket requests:
  - <https://rt.hpc.msu.edu/index.html>
  - Please include "Bioinformatics Help" in the subject to more quickly route your request.

MICHIGAN STATE UNIVERSITY

ICER

## Doing more faster

Pleasantly Parallel Techniques

MICHIGAN STATE UNIVERSITY

ICER

## How much Communication?

- Pleasantly parallel
  - No communication required
- Loosely coupled
  - Typically sync at regular intervals
- Tightly coupled
  - Constant communication

MICHIGAN STATE UNIVERSITY

ICER

## Pleasantly Parallel

MICHIGAN STATE UNIVERSITY

ICER

## How fast can we go?

- T - How long does each operation take?
- N - How many operations do you need to run?
- CPUs – Number of Cores job will run on.
- Single CPU time estimate:
  - $T \times N$
- Best possible Pleasantly parallel time:
  - $(T \times N) * \text{overhead} / \text{CPUs}$

MICHIGAN STATE UNIVERSITY

ICER

## Who are you? -- Biometrics

MICHIGAN STATE UNIVERSITY

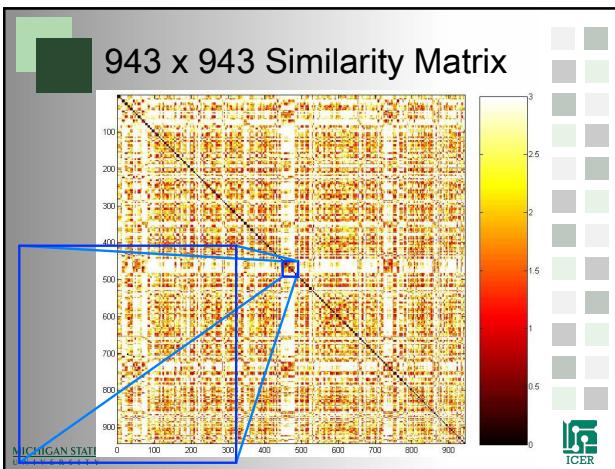
ICER

## Pairwise-All Problem

- Database of faces
- Compare everything to everything else
- Calculate a Matching score to use for identification <http://tech.msu.edu/CI-Days/>

MICHIGAN STATE UNIVERSITY

ICER



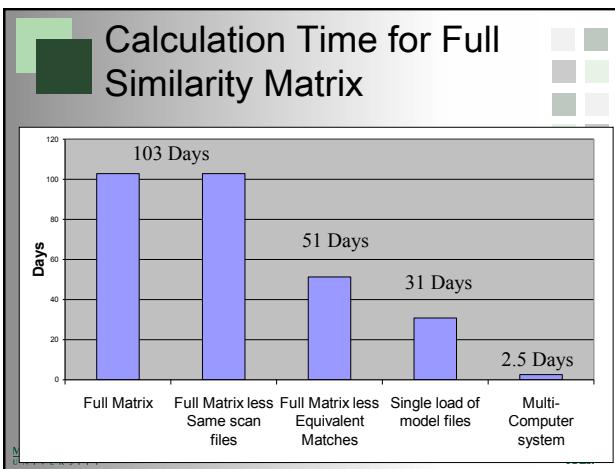
## Estimated Calculation Times

- Preprocessing
  - $943 * 12$  (seconds)  $\approx$  189 Minutes
- Matching
  - $943 * 943 * 5$  (seconds)  $\approx$  103 Days
  - Scans matched to themselves always result in 0 mm
    - $(943 * 943 - 943) * 5$  (seconds)  $\approx$  103 Days
- The Proposed Alignment Algorithm is symmetric.
  - $(943 * 943 - 943)/2 * 5$  (seconds)  $\approx$  51.5 Days
- We also load models once per row instead of every time
  - $(943*943-943)/2 * 3$  (seconds) +  $943 * 2$  (seconds)  $\approx$  31 Days

40

MICHIGAN STATE UNIVERSITY

ICER



## How do we go even bigger?

- 5000 scans.
  - 1.5 years on a single processor computer
  - 13 days on our ad-hoc cluster.
  - 1.5 days on the HPCC

42

MICHIGAN STATE UNIVERSITY

ICER

## Step to Pleasantly Parallel

- Figure out command line
- Estimate single job time:
  - Should be > 5 minutes
  - Should be < 1 week
  - Best if < 4 hours
- Make a submissions script
- Submit Job

MICHIGAN STATE  
UNIVERSITY

## PBS Job Arrays

- One submission script copied many times
- Uses the PBS -t option
  - Ranges: 1-10
  - Lists: 2,4,100,3
  - Combination: 1-10,20,50,100
- Distinguish between jobs by using the PBS\_ARRAYID environment variable

MICHIGAN STATE  
UNIVERSITY

## Example: Job Arrays

- Get the blender\_farm example:
 

```
>module load powertools
>getexample blender_farm
>cd ./blender_farm
```
- Look at the qsub file, using "less" command
 

```
>less blender_farm.qsub
```
- Submit the job
 

```
>qsub blender_farm.qsub
```

MICHIGAN STATE  
UNIVERSITY

## Example

- Folder full of input files:

1.in	5.in	9.in	13.in	17.in
2.in	6.in	10.in	14.in	18.in
3.in	7.in	11.in	15.in	19.in
4.in	8.in	12.in	16.in	

- Want folder full of output files:

1.out	5.out	9.out	13.out	17.out
2.out	6.out	10.out	14.out	18.out
3.out	7.out	11.out	15.out	19.out
4.out	8.out	12.out	16.out	

- Command Syntax:

- ./myprogram inputfile > outputfile

MICHIGAN STATE  
UNIVERSITY

## Simple Job Array

```
#!/bin/bash -login
#PBS -l walltime=00:05:00,mem=2gb
#PBS -l nodes=1:ppn=1,feature=gbe
#PBS -t 1-100

cd ${PBS_O_WORKDIR}

./myprogram ${PBS_ARRAYID}.in > ${PBS_ARRAYID}.out

qstat -f ${PBS_JOBID}
```

MICHIGAN STATE  
UNIVERSITY

## HPCC Job array limitations

- Can not have more than 144 cores running
- Can not submit more than 256 jobs at once
- Lots of ways to work around this problem.

MICHIGAN STATE  
UNIVERSITY

 Job array numbers

- All numbers in a job array have the same base number
  - 7478210
- Each PBS\_ARRAYID is shown in square brackets
  - 7478210[1]
  - 7478210[2]
- Delete all jobs using one command
  - qdel 7478210[]





 **Condor**  
High Throughput Computing

- Job submission system
- Runs like a screen saver
- Steals CPU Cycles from computers on campus







 Bioinformatics Contact

- John Johnston, HPC Programmer
  - M-W, 1449 BPS, 884-2572
  - Th-F, 505 BMB, 432-7177
  - [johnj@msu.edu](mailto:johnj@msu.edu)
- Ticket requests:
  - <https://rt.hpc.msu.edu/index.html>
  - Please include “Bioinformatics Help” in the subject to more quickly route your request.



